

A dark blue vertical bar runs down the left side of the page. A blue arrow points to the right from the bar, containing the year 2024.

2024

Insurance Premium Prediction

Capstone Project - Regression

Several thin, curved lines in dark blue and light grey originate from the bottom left corner and sweep upwards and to the right.

DIBOSH BARUAH

1. Introduction

Health insurance premium prediction is a crucial task in the insurance industry as it helps both insurance companies and policyholders make informed decisions. Predicting accurate premiums is essential for risk assessment and ensuring that policyholders pay premiums that reflect their health risks. In this project, the goal is to predict health insurance premiums in the United States based on various personal and health-related factors, such as age, BMI, smoking status, medical history, and others. A **Linear Regression model** was chosen due to its simplicity and interpretability in regression tasks. The dataset includes 10 variables related to the individual's demographics and health behaviors, with the target variable being the insurance premium charges.

2. Data Collection

The dataset used for this project contains various attributes influencing the cost of health insurance premiums in the U.S. The features include:

- **Age:** The age of the individual.
- **Gender:** The gender of the individual (Male/Female).
- **BMI (Body Mass Index):** A measure of body fat based on height and weight.
- **Children:** The number of children/dependents covered by the insurance.
- **Smoker:** Whether the individual is a smoker (Yes/No).
- **Region:** The region in which the individual resides (northeast, northwest, southeast, southwest).
- **Medical History:** The individual's medical history (e.g., diabetes, hypertension, etc.).
- **Family Medical History:** The family history of medical conditions (e.g., heart disease, diabetes).
- **Exercise Frequency:** How often the individual exercises (Never, Occasionally, Rarely).
- **Occupation:** The individual's job type (e.g., White collar, Blue collar).
- **Coverage Level:** The type of insurance coverage (Standard, Premium).
- **Charges:** The target variable representing the health insurance premium amount.

The dataset was loaded into a Pandas DataFrame, and preprocessing steps were carried out to clean and prepare the data for modeling.

3. Data Preprocessing

Several preprocessing steps were performed to prepare the data for the Linear Regression model:

3.1 Handling Missing Values

The dataset was checked for missing values. Any rows containing missing values were filled with 0, and no significant missing values were found in the columns after preprocessing.

3.2 Data Type Conversion

Columns were converted to appropriate data types:

- **Numeric Columns** (e.g., age, bmi, children, charges) were converted to float or int.
- **Categorical Columns** (e.g., gender, smoker, region, medical history) were converted to category data type to optimize memory usage and processing.

3.3 One-Hot Encoding

Categorical variables were transformed into numerical values using one-hot encoding. This approach helps in converting non-numeric categories into binary format suitable for machine learning algorithms. The first category of each feature was dropped to avoid multicollinearity.

3.4 Feature Engineering

- A new feature, **BMI Category**, was created by categorizing BMI values into bins (Underweight, Normal, Overweight, Obese). This transformation was useful for better representing the BMI variable for prediction.
- **One-Hot Encoding** was applied to the newly created BMI category to integrate it with the rest of the features.

3.5 Feature Scaling

No feature scaling was applied as Linear Regression is generally not sensitive to feature scales. However, it's a good practice to standardize features when using more complex models.

4. Model Building

The goal was to build a predictive model that estimates health insurance premiums based on the input features. The steps involved in building the model are as follows:

4.1 Splitting the Dataset

The dataset was split into **features** (x) and **target variable** (y, which is the charges/insurance premium). The data was then split into training and testing sets using an 80-20 ratio.

4.2 Training the Model

A **Linear Regression model** was trained on the training data. Linear regression was chosen for its simplicity and interpretability, as it provides insights into how each feature influences the target variable (charges).

4.3 Making Predictions

Once the model was trained, it was used to make predictions on the test set. The predicted premiums were then compared to the actual premium values to evaluate the model's performance.

5. Model Evaluation

The model's performance was evaluated using several key metrics to assess its predictive accuracy:

- **Mean Squared Error (MSE)**: A measure of the average squared difference between the predicted and actual values.
- **Mean Absolute Error (MAE)**: The average of the absolute differences between the predicted and actual values.
- **Root Mean Squared Error (RMSE)**: The square root of the MSE, providing a more interpretable measure of prediction error.
- **R-Squared (R^2)**: A metric indicating the proportion of variance in the target variable (charges) that is explained by the features.

5.1 Performance Metrics

The evaluation metrics on the test set were as follows:

- **Mean Squared Error (MSE):** 136,042.26
- **R-Squared (R^2):** 0.9930
- **Mean Absolute Error (MAE):** 292.86
- **Root Mean Squared Error (RMSE):** 368.84

These metrics indicate that the model is performing quite well, explaining over 99% of the variance in the target variable, with a relatively low error rate.

6. Results & Discussion

6.1 Model Performance

The **Linear Regression model** performed excellently with an **R^2 value of 0.9930**, meaning that 99.3% of the variation in insurance premiums was explained by the input features. The **RMSE** of 368.84 suggests that the model's predictions are reasonably close to the actual charges, with a slight deviation. However, there's always room for improvement.

6.2 Feature Importance

Linear regression provides insight into the importance of each feature through its **coefficients**. The top features contributing to the prediction of health insurance premiums were:

- **Age:** Older individuals generally pay higher premiums.
- **BMI:** Higher BMI values lead to higher premiums due to associated health risks.
- **Smoker Status:** Smokers pay higher premiums compared to non-smokers due to higher health risks.
- **Medical History:** Having pre-existing medical conditions like diabetes increases the premium.
- **Number of Children:** Individuals with more children tend to have higher premiums due to higher family healthcare costs.

These features significantly influenced the predicted premiums, and understanding their relationships with the target variable can help optimize insurance policies.

7. Conclusion

This project successfully demonstrated the use of **Linear Regression** to predict health insurance premiums in the U.S. The model performed well with an **R^2 of 0.9930**, indicating that it could explain most of the variance in insurance premiums. The key features influencing the predictions were age, BMI, smoking status, and medical history. Further improvements could be made by exploring more advanced algorithms like **Random Forests** or **Gradient Boosting**, or by tuning hyperparameters for better accuracy.

8. Original Code:

Original coding process

```
import sqlite3
import pandas as pd
```

Connecting to the database

```
conn = sqlite3.connect('/Users/diboshbaruah/Desktop/Database.db')
data = pd.read_sql_query('SELECT * FROM Insurance_Prediction', conn)
```

```
print("Dataset successfully loaded...\n")
```

Displaying the first few rows to inspect the data

```
print("Displaying first few rows of the dataset:\n")
print(data.head())
```

Checking data types before conversion

```
print("\nData types before conversion:")
print(data.dtypes)
```

Closing the connection

```
conn.close()
```

Dataset successfully loaded...

Displaying first few rows of the dataset:

	age	gender	bmi	children	smoker	region	medical_history \
0	46.0	male	21.45	5.0	yes	southeast	Diabetes
1	25.0	female	25.38	2.0	yes	northwest	Diabetes
2	38.0	male	44.88	2.0	yes	southwest	
3	25.0	male	19.89	0.0	no	northwest	
4	49.0	male	38.21	3.0	yes	northwest	Diabetes

	family_medical_history	exercise_frequency	occupation	coverage_level \
0		Never	Blue collar	Premium
1	High blood pressure	Occasionally	White collar	Premium
2	High blood pressure	Occasionally	Blue collar	Premium
3	Diabetes	Rarely	White collar	Standard
4	High blood pressure	Rarely	White collar	Standard

	charges
0	20460.307668871566
1	20390.8992176422
2	20204.476301934814
3	11789.029842697417
4	19268.309838159606

Data types before conversion:

```
age                object
gender             object
bmi               object
children          object
smoker            object
region            object
medical_history    object
family_medical_history object
exercise_frequency object
occupation         object
coverage_level     object
charges           object
dtype: object
```

Converting columns to numeric where applicable

```
numeric_cols = ['age', 'bmi', 'children', 'charges']
data[numeric_cols] = data[numeric_cols].apply(pd.to_numeric, errors='coerce')
```

Calculating Q1, Q3, and IQR for each numerical column

```
Q1 = data[numeric_cols].quantile(0.25)
Q3 = data[numeric_cols].quantile(0.75)
IQR = Q3 - Q1
```

Defining the outlier boundaries

```
lower_bound = Q1 - 1.5 * IQR
upper_bound = Q3 + 1.5 * IQR
```

Identifying outliers by applying the boundaries to the data

```
outliers = ((data[numeric_cols] < lower_bound) | (data[numeric_cols] > upper_bound))
).any(axis=1)
```

```
# Counting the number of outliers
num_outliers = outliers.sum()
print(f"Number of outliers: {num_outliers}")
```

```
# Displaying the rows containing outliers
outlier_rows = data[outliers]
print("\nRows containing outliers:")
print(outlier_rows)
```

Number of outliers: 2217

Rows containing outliers:

	age	gender	bmi	children	smoker	region	medical_history \
59	49.0	female	44.67	5.0	yes	northwest	Heart disease
117	46.0	male	42.33	2.0	yes	southwest	Heart disease
704	53.0	female	48.23	3.0	yes	southeast	Heart disease
804	53.0	female	35.30	0.0	yes	southwest	Heart disease
1451	63.0	male	41.76	2.0	yes	northwest	Heart disease
...
998316	NaN	male	39.52	3.0	yes	northeast	Heart disease
998667	65.0	male	49.77	4.0	yes	southeast	Heart disease
999077	37.0	female	22.70	0.0	yes	northeast	Heart disease
999319	25.0	female	42.41	4.0	yes	northeast	Heart disease
999790	64.0	male	44.41	3.0	yes	southeast	Heart disease

	family_medical_history	exercise_frequency	occupation	coverage_level \
59	Heart disease	Frequently	White collar	Premium
117	Heart disease	Frequently	Student	Premium
704	Heart disease	Occasionally	White collar	Premium
804	Heart disease	Frequently	White collar	Premium
1451	Heart disease	Rarely	White collar	Premium
...
998316	Heart disease	Frequently	Student	Premium
998667	Heart disease	Rarely	White collar	Premium
999077	Heart disease	Frequently	White collar	Premium
999319	Heart disease	Frequently	White collar	Premium
999790	Heart disease	Occasionally	Blue collar	Premium

	charges
59	30370.731957
117	29412.861604
704	29913.058400
804	29183.845771
1451	29693.427087
...	...
998316	30017.590027
998667	30628.879005
999077	29204.562471
999319	30566.944707
999790	30074.556653

[2217 rows x 12 columns]


```

import matplotlib.pyplot as plt
import seaborn as sns

# Box plots for outliers
plt.figure(figsize=(12, 8))

# Box plot for age
plt.subplot(2, 2, 1)
sns.boxplot(x=data['age'])
plt.title('Distribution of Age')

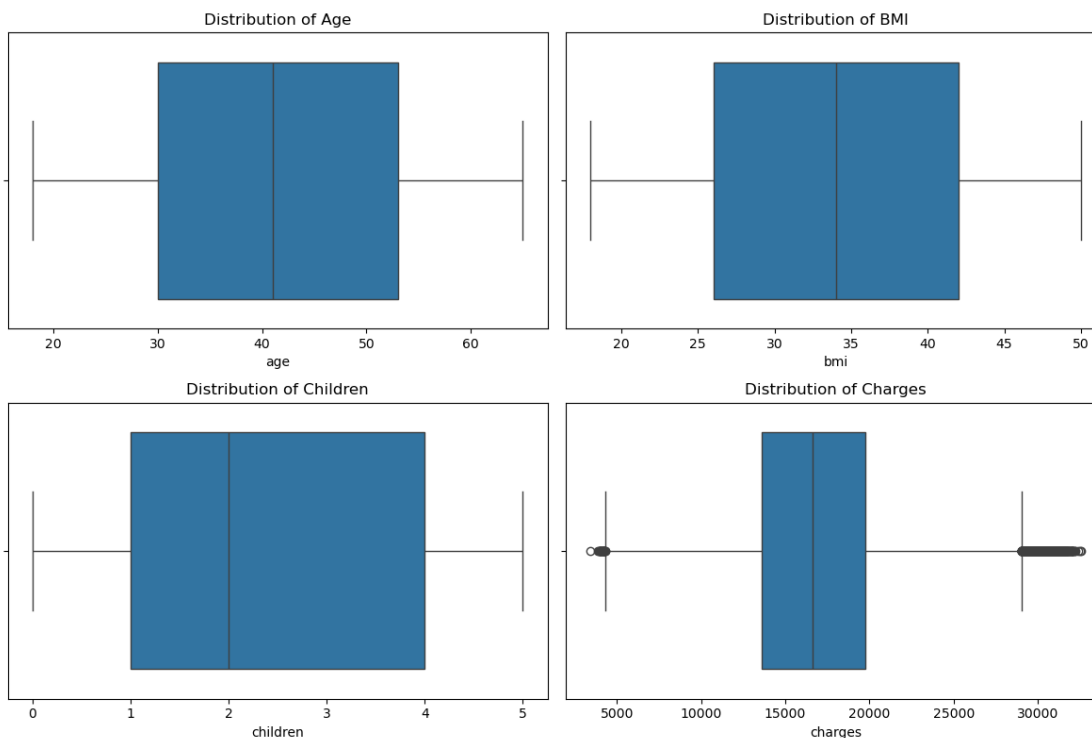
# Box plot for BMI
plt.subplot(2, 2, 2)
sns.boxplot(x=data['bmi'])
plt.title('Distribution of BMI')

# Box plot for children
plt.subplot(2, 2, 3)
sns.boxplot(x=data['children'])
plt.title('Distribution of Children')

# Box plot for charges
plt.subplot(2, 2, 4)
sns.boxplot(x=data['charges'])
plt.title('Distribution of Charges')

# Adjust Layout
plt.tight_layout()
plt.show()

```



```
# Data Pre-processing
print("Initiating data pre-processing..")

# Converting columns to appropriate data types (if not already done)
data['age'] = pd.to_numeric(data['age'], errors='coerce')
data['bmi'] = pd.to_numeric(data['bmi'], errors='coerce')
data['children'] = pd.to_numeric(data['children'], errors='coerce')
data['charges'] = pd.to_numeric(data['charges'], errors='coerce')

# Converting categorical columns to 'category' dtype
data['gender'] = data['gender'].astype('category')
data['smoker'] = data['smoker'].astype('category')
data['region'] = data['region'].astype('category')
data['medical_history'] = data['medical_history'].astype('category')
data['family_medical_history'] = data['family_medical_history'].astype('category')
data['exercise_frequency'] = data['exercise_frequency'].astype('category')
data['occupation'] = data['occupation'].astype('category')
data['coverage_level'] = data['coverage_level'].astype('category')

# Performing one-hot encoding on the categorical columns
data_encoded = pd.get_dummies(data, columns=[
    'gender', 'smoker', 'region', 'medical_history',
    'family_medical_history', 'exercise_frequency', 'occupation', 'coverage_level']
, drop_first=True)

# Filling NaN values with 0
data_encoded = data_encoded.fillna(0)

# Converting boolean columns (True/False) to integer (1/0)
data_encoded = data_encoded.astype(int)
```

Initiating data pre-processing..

```
# Checking data types after conversion
```

```
print(data_encoded.head())
```

	age	bmi	children	charges	gender_female	gender_male	smoker_yes	\
0	46	21	5	20460	0	1	1	
1	25	25	2	20390	1	0	1	
2	38	44	2	20204	0	1	1	
3	25	19	0	11789	0	1	0	
4	49	38	3	19268	0	1	1	

	region_northwest	region_southeast	region_southwest	...	\
0	0	1	0	...	
1	1	0	0	...	
2	0	0	1	...	
3	1	0	0	...	
4	1	0	0	...	

	family_medical_history_High blood pressure	exercise_frequency_Never	\
0	0	1	
1	1	0	
2	1	0	
3	0	0	
4	1	0	

	exercise_frequency_Occasionally	exercise_frequency_Rarely	\
0	0	0	
1	1	0	
2	1	0	
3	0	1	
4	0	1	

	occupation_Blue collar	occupation_Student	occupation_Unemployed	\
0	1	0	0	
1	0	0	0	
2	1	0	0	
3	0	0	0	
4	0	0	0	

	occupation_White collar	coverage_level_Premium	coverage_level_Standard
0	0	1	0
1	1	1	0
2	0	1	0
3	1	0	1
4	1	0	1

```
[5 rows x 25 columns]
```

```
# Checking for missing values in the dataset after conversion
```

```
missing_values = data_encoded.isnull().sum()  
print("Missing values for each column:")  
print(missing_values)
```

```
Missing values for each column:
```

```
age                                0  
bmi                                0  
children                          0  
charges                           0  
gender_female                     0  
gender_male                       0  
smoker_yes                        0  
region_northwest                  0  
region_southeast                  0  
region_southwest                  0  
medical_history_Diabetes          0  
medical_history_Heart disease    0  
medical_history_High blood pressure 0  
family_medical_history_Diabetes  0  
family_medical_history_Heart disease 0  
family_medical_history_High blood pressure 0  
exercise_frequency_Never         0  
exercise_frequency_Occasionally  0  
exercise_frequency_Rarely        0  
occupation_Blue collar           0  
occupation_Student               0  
occupation_Unemployed            0  
occupation_White collar          0  
coverage_level_Premium           0  
coverage_level_Standard          0  
dtype: int64
```

```
# Feature Engineering
```

```
bins = [0, 18.5, 24.9, 29.9, float('inf')]  
labels = ['Underweight', 'Normal', 'Overweight', 'Obese']  
data_encoded['bmi_category'] = pd.cut(data_encoded['bmi'], bins=bins, labels=labels  
)
```

```
# One-hot encoding bmi_category
```

```
data_encoded = pd.get_dummies(data_encoded, columns=['bmi_category'], drop_first=True  
ue)
```

```
# Splitting the dataset into training, evaluation, and live data
```

```
# Splitting first 700k records for training (train size = 700,000)
```

```
train_data = data_encoded.iloc[:700000]  
X_train = train_data.drop(columns=['charges']) # Features  
y_train = train_data['charges'] # Target variable
```

```
# Splitting next 200k records for evaluation (eval size = 200,000)
```

```
eval_data = data_encoded.iloc[700000:900000]  
X_eval = eval_data.drop(columns=['charges']) # Features  
y_eval = eval_data['charges'] # Target variable
```

```

# Will use remaining 100k records as live data (Live size = 100,000)
live_data = data_encoded.iloc[900000:]
X_live = live_data.drop(columns=['charges']) # Features
y_live = live_data['charges'] # Target variable

from sklearn.ensemble import RandomForestRegressor

# Initializing the Random Forest Regressor
rf_regressor = RandomForestRegressor(n_estimators=20, random_state=42)

# Training the model on the training data
rf_regressor.fit(X_train, y_train)

RandomForestRegressor(n_estimators=20, random_state=42)

from sklearn.metrics import mean_squared_error, r2_score, mean_absolute_error
import numpy as np

# Model Prediction and Evaluation on evaluation dataset
print("\nModel Evaluation on the evaluation dataset...\n")

# Predictions on evaluation set
y_eval_pred = rf_regressor.predict(X_eval)

# Evaluating the model on evaluation dataset
mse = mean_squared_error(y_eval, y_eval_pred)
r2 = r2_score(y_eval, y_eval_pred)
mae = mean_absolute_error(y_eval, y_eval_pred)
rmse = np.sqrt(mse)

print(f"Mean Squared Error (Evaluation Set): {mse}")
print(f"R^2 (Evaluation Set): {r2}")
print(f"Mean Absolute Error (Evaluation Set): {mae}")
print(f"Root Mean Squared Error (Evaluation Set): {rmse}")

# Model Evaluation on live dataset
print("\nModel Evaluation on live dataset...\n")

# Predictions on live data
y_live_pred = rf_regressor.predict(X_live)

# Evaluating the model on live dataset
mse_live = mean_squared_error(y_live, y_live_pred)
r2_live = r2_score(y_live, y_live_pred)
mae_live = mean_absolute_error(y_live, y_live_pred)
rmse_live = np.sqrt(mse_live)

print(f"Mean Squared Error (Live Set): {mse_live}")
print(f"R^2 (Live Set): {r2_live}")
print(f"Mean Absolute Error (Live Set): {mae_live}")
print(f"Root Mean Squared Error (Live Set): {rmse_live}")

```

Model Evaluation on the evaluation dataset...

Mean Squared Error (Evaluation Set): 222379.48812875623
R^2 (Evaluation Set): 0.9886071477464806
Mean Absolute Error (Evaluation Set): 370.4485876982141
Root Mean Squared Error (Evaluation Set): 471.5712969729564

Model Evaluation on live dataset...

Mean Squared Error (Live Set): 221656.89872781327
R^2 (Live Set): 0.9885407292677871
Mean Absolute Error (Live Set): 370.87286679583343
Root Mean Squared Error (Live Set): 470.8045228412884

** Now running the saved scripts on jupyter notebook - train_model.py // predict_fraud.py // app.py

Importing the training script
!python train_model.py

Model trained and saved successfully!

Importing the Predict script
!python predict_model.py

Model Evaluation on the evaluation dataset...

Mean Squared Error (Evaluation Set): 222379.48812875623
R^2 (Evaluation Set): 0.9886071477464806
Mean Absolute Error (Evaluation Set): 370.4485876982141
Root Mean Squared Error (Evaluation Set): 471.5712969729564

Model Evaluation on live dataset...

Mean Squared Error (Live Set): 221656.89872781327
R^2 (Live Set): 0.9885407292677871
Mean Absolute Error (Live Set): 370.87286679583343
Root Mean Squared Error (Live Set): 470.8045228412884

import subprocess

Now running the Flask app using subprocess
subprocess.Popen(["python", "app.py"])

<Popen: returncode: None args: ['python', 'app.py']>

```
# Sample input data
```

```
import requests
```

```
data = {  
    'age': 60,  
    'bmi': 28.5,  
    'children': 1,  
    'gender_male': 1,  
    'smoker_yes': 0,  
    'region_northwest': 0,  
    'region_southeast': 1,  
    'region_southwest': 0,  
    'medical_history_yes': 1,  
    'family_medical_history_yes': 0,  
    'exercise_frequency_high': 0,  
    'exercise_frequency_medium': 1,  
    'exercise_frequency_low': 0,  
    'occupation_occupation1': 0,  
    'occupation_occupation2': 1,  
    'coverage_level_high': 0,  
    'coverage_level_low': 1,  
    'bmi_category_Overweight': 1,  
    'bmi_category_Obese': 0  
}
```

```
# API endpoint
```

```
url = 'http://127.0.0.1:5000/predict'
```

```
# Sending POST request
```

```
response = requests.post(url, json=data)
```

```
# Printing the result
```

```
try:  
    print(response.json())  
except Exception as e:  
    print(f"Error in response: {e}")  
    print(f"Response content: {response.text}")
```

```
{'prediction': 9437.538769246914}
```